



Improving scheduling latency with SCHED_IDLE task

Viresh Kumar (Linaro)



Linaro
connect
San Diego 2019

Scheduling classes and policies (recap)

- Stop
 - No policy
- Deadline
 - SCHED_DEADLINE
- Real Time
 - SCHED_FIFO
 - SCHED_RR
- Fair (CFS)
 - SCHED_NORMAL
 - SCHED_BATCH
 - SCHED_IDLE
- Idle
 - No policy

SCHED_IDLE scheduling policy

- One of the scheduling policies available in CFS.
- Nice range -20 to +19, where +19 means lowest priority.
- Priority of SCHED_IDLE tasks is lower than tasks with +19 nice.
- Useful for background tasks with lowest priority in system.
- Any non-idle task can preempt a running SCHED_IDLE task.
- It is still not a real idle time scheduler and still gets a chance to run on a fully busy system to avoid priority inversion.
- Not widely used currently. That needs to change.

What's new ?

Before:

- Special handling only in `->check_preempt_curr()` in `fair.c`
 - A `SCHED_OTHER` task will preempt a `SCHED_IDLE` task.

Now (Merged in 5.4-rc1):

- Special handling added in `->select_task_rq()` in `fair.c`
 - Both fast and slow paths updated
 - CPU running only `SCHED_IDLE` tasks considered as idle
 - Will immediately get preempted by a `SCHED_NORMAL` task.
 - Improves scheduling latency for the `SCHED_OTHER` task
 - Even better than running on an idle-cpu as we don't need to wake-up the CPU.

rt-app json, 8 normal and 5 sched-idle threads

```
{
  "tasks": {
    "cfs_thread": {
      "instance": 8,
      "run": 5333,
      "timer": { "ref": "unique", "period": 7777 },
      "policy": "SCHED_OTHER"
    },
    "idle_thread": {
      "instance": 5,
      "run": 3000,
      "policy": "SCHED_IDLE"
    }
  },
  "global": {
    "duration": 5,
    "calibration": "CPU0",
    "default_policy": "SCHED_OTHER",
    "pi_enabled": false,
    (more unrelated settings...)
    "gnuplot": false
  }
}
```

Wu-latency (usec) from rt-app for CFS tasks

Hardware: Octa core cortex A7

Without the patchset:


N	min	max	sum	mean	stddev
4710	0	67664	5.25956e+06	1116.68	2315.09

With the patchset:


N	min	max	sum	mean	stddev
5095	0	7773	523170	102.683	475.482

mean latency dropped by 90% and stddev dropped by 75%

Kernelshark without sched-idle modifications



Kernelshark without sched-idle modifications ...




The image displays a CPU timeline visualization for 8 CPUs (CPU 0 to CPU 7). The timeline shows the execution of various threads, with labels indicating the thread name and its ID. The threads include:

- idle_thread-8-622**: Running on CPU 0.
- cfs_thread-3-617**: Running on CPU 2.
- cfs_thread-1-615**: Running on CPU 3.
- cfs_thread-4-618**: Running on CPU 4.
- cfs_thread-7-621**: Running on CPU 4.
- cfs_thread-2-616**: Running on CPU 4.
- idle_thread-9-623**: Running on CPU 6.


The timeline is divided into segments for different threads, with labels indicating the thread name and its ID. The threads include:

- idle_thread-8-622**: Running on CPU 0.
- cfs_thread-3-617**: Running on CPU 2.
- cfs_thread-1-615**: Running on CPU 3.
- cfs_thread-4-618**: Running on CPU 4.
- cfs_thread-7-621**: Running on CPU 4.
- cfs_thread-2-616**: Running on CPU 4.
- idle_thread-9-623**: Running on CPU 6.

Kernelshark with sched-idle modifications



Kernelshark with sched-idle modifications



Who should use SCHED_IDLE policy ?

- SCHED_IDLE policy isn't widely used currently.
- Can be used for most of background tasks which aren't time critical.
- Google is interested in using SCHED_IDLE policy for background Android tasks, like dex2oat (compiles dex files).
- Facebook's use-case involves using spare CPU cycles on servers (running latency sensitive workloads) to run side jobs like video encoding but that interferes with main workload. Making the side jobs SCHED_IDLE has proven to be very useful.

Thank you

Join Linaro to accelerate deployment of your
Arm-based solutions through collaboration

contactus@linaro.org



96boards is a range of specifications with
boards and peripherals offering different
performance levels and features in a
standard footprint.



Linaro
connect

San Diego 2019